

Problem 1: In this problem set, we are going to be considering a variant on the hierarchical bag-of-words model that we looked at in class. In class, we used a Dirichlet distribution to define a prior distribution over θ , the parameters of the bag of words model. The Dirichlet distribution is a continuous distribution on the simplex—it assigns probability mass to all the uncountably many points on the simplex.

For this problem set, we will be looking at a considerably simpler prior distribution over the parameters θ . Our distribution will be *discrete*, and in particular will only assign positive probability to a finite number of values of θ . The probability distribution is defined in the code below:

```
(def vocabulary '(call me ishmael))

(def theta1 (list (/ 1 2) (/ 1 4) (/ 1 4)))
(def theta2 (list (/ 1 4) (/ 1 2) (/ 1 4)))

(def thetas (list theta1 theta2))

(def theta-prior (list (/ 1 2) (/ 1 2)))
```

Our vocabulary in this case consists of three words. Each value of θ therefore defines a bag of words distribution over sentences containing these three words. The first value of θ (`theta1`) assigns $\frac{1}{2}$ probability to the word ‘call’, $\frac{1}{4}$ to ‘me’, and $\frac{1}{4}$ to ‘ishmael’. The second value of θ (`theta2`) assigns $\frac{1}{2}$ probability to ‘me’, and $\frac{1}{4}$ to each of the other two words. The two values of θ each have prior probability of $\frac{1}{2}$. *Assume throughout the problem set that the vocabulary and possible values of θ are fixed to these values above.*

In addition to the code defining the prior distribution over θ , we will be using some helper functions defined in class:

```
(defn score-categorical [outcome outcomes params]
  (if (empty? params)
      (throw "no matching outcome")
      (if (= outcome (first outcomes))
          (first params)
          (score-categorical outcome (rest outcomes) (rest params)))))

(defn list-foldr [f base lst]
  (if (empty? lst)
      base
      (f (first lst)
         (list-foldr f base (rest lst)))))

(defn score-BOW-sentence [sen probabilities]
  (list-foldr
   (fn [word rest-score]
     (+ (log2 (score-categorical word vocabulary probabilities))
```

```
        rest-score))
    0
    sen))

(defn score-corpus [corpus probabilities]
  (list-foldr
    (fn [sen rst]
      (+ (score-BOW-sentence sen probabilities) rst))
    0
    corpus))

(defn logsumexp [log-vals]
  (let [mx (apply max log-vals)]
    (+ mx
      (log2
        (apply +
          (map (fn [z] (Math/pow 2 z))
                (map (fn [x] (- x mx)) log-vals)))))))
```

Recall that the function `score-corpus` is used to compute the log probability of a corpus given a particular value of the parameters θ . Also recall (from the Discrete Random Variables module) the purpose of the function `logsumexp`, which is used to compute the sum of (log) probabilities; you should return to the lecture notes if you don't remember what this function is doing. (Note that the version of `logsumexp` here differs slightly from the lecture notes, as it does not use the `&` notation.)

Our initial corpus will consist of two sentences:

```
(def my-corpus '((call me)
                  (call ishmael)))
```

Write a function `theta-corpus-joint`, which takes three arguments: `theta`, `corpus`, and `theta-probs`. The argument `theta` is a value of the model parameters θ , and the argument `corpus` is a list of sentences. The argument `theta-probs` is a prior probability distribution over the values of θ . The function should return the **log** of the joint probability $\Pr(C = \text{corpus}, \theta = \text{theta})$.

Use the chain-rule identity discussed in class: $\Pr(C, \theta) = \Pr(C|\theta) \Pr(\theta)$. Assume that the prior distribution $\Pr(\theta)$ is defined by the probabilities in `theta-probs`.

After defining this function, call `(theta-corpus-joint theta1 my-corpus theta-prior)`. This will compute (the log of) the joint probability of the model parameters `theta1` and the corpus `my-corpus`.

Answer 1: <https://snipplr.com/view/65234/logarithm-with-base-2-log2/>

```
(defn log2 [n] (/ (Math/log n) (Math/log 2)))
```

```
(defn theta-corporus-joint [theta corpus theta-probs]
  (if (= theta (first thetas))
    (+ (score-corporus corpus theta)(log2 (first theta-probs)))
    (+ (score-corporus corpus theta)(log2 (first (rest theta-probs)))))))
```

Problem 2: Write a function `compute-marginal`, which takes two arguments: `corpus` and `theta-probs`. The argument `corpus` is a list of sentences, and the argument `theta-probs` is a prior probability distribution on values of θ . The function should return the **log** of the marginal likelihood of the corpus, when the prior distribution on θ is given by `theta-probs`. That is, the function should return $\log[\sum_{\theta \in \Theta} \Pr(\mathbf{C} = \text{corpus}, \Theta = \theta)]$.

Hint: Use the `logsumexp` function defined above.

After defining `compute-marginal`, call `(compute-marginal my-corpus theta-prior)`. This will compute the marginal likelihood of `my-corpus` (which was defined above), given the prior distribution `theta-prior`.

Answer 2:

```
(defn compute-marginal [corpus theta-probs]
  (logsumexp (list (theta-corporus-joint theta1 corpus theta-probs)
                  (theta-corporus-joint theta2 corpus theta-probs))))
(compute-marginal my-corpus theta-prior)
-6.415037499278844
(Math/pow 2 (compute-marginal my-corpus theta-prior))
0.01171875
```

Problem 3: Write a procedure `compute-conditional-prob`, which takes three arguments: `theta`, `corpus`, and `theta-probs`. The arguments have the same interpretation as in Problems 1 and 2. The function should return the **log** of the conditional probability of the parameter value `theta`, given the corpus. Remember that the conditional probability is defined by the equation:

$$\Pr(\Theta = \theta | \mathbf{C} = \text{corpus}) = \frac{\Pr(\mathbf{C} = \text{corpus}, \Theta = \theta)}{\sum_{\theta \in \Theta} \Pr(\mathbf{C} = \text{corpus}, \Theta = \theta)} \quad (1)$$

Answer 3:

$$\begin{aligned}\Pr(\Theta = \theta | \mathbf{C} = \text{corpus}) &= \frac{\Pr(\mathbf{C} = \text{corpus}, \Theta = \theta)}{\sum_{\theta \in \Theta} \Pr(\mathbf{C} = \text{corpus}, \Theta = \theta)} \\ \log_2(\Pr(\Theta = \theta | \mathbf{C} = \text{corpus})) &= \log_2 \left(\frac{\Pr(\mathbf{C} = \text{corpus}, \Theta = \theta)}{\sum_{\theta \in \Theta} \Pr(\mathbf{C} = \text{corpus}, \Theta = \theta)} \right) \\ &= \log_2(\Pr(\mathbf{C} = \text{corpus}, \Theta = \theta)) - \log_2 \left(\sum_{\theta \in \Theta} \Pr(\mathbf{C} = \text{corpus}, \Theta = \theta) \right)\end{aligned}$$

```
(defn compute-conditional-prob [theta corpus theta-probs]
  (- (theta-corpus-joint theta corpus theta-probs)
     (compute-marginal corpus theta-probs)))
(compute-conditional-prob theta1 my-corpus theta-prior)
-0.5849625007211561
```

Problem 4: Write a function `compute-conditional-dist`, which takes two arguments: `corpus` and `theta-probs`. For every value of θ in `thetas` (i.e., `theta1` and `theta2`), it should return the conditional probability of θ given the corpus. That is, it should return a list of conditional probabilities of the different values of θ .

Answer 4:

```
(defn compute-conditional-dist [corpus theta-probs]
  (list (compute-conditional-prob theta1 corpus theta-probs)
        (compute-conditional-prob theta2 corpus theta-probs)))
```

Problem 5: Call `(compute-conditional-dist my-corpus theta-prior)`. What do you notice about the conditional distribution over values of θ ? Exponentiate the values you get back, so that you can see the regular probabilities, rather than just the log probabilities. Explain why the conditional distribution looks the way it does, with reference to the properties of `my-corpus`.

Answer 5:

```
(compute-conditional-dist my-corpus theta-prior)
(-0.5849625007211561 -1.584962500721156)
(defn exponentiate [lst]
  (if (empty? lst)
```

```
'()
  (cons (Math/pow 2 (first lst)) (exponentiate (rest lst))))
(exponentiate (compute-conditional-dist my-corpus theta-prior))
(0.6666666666666667 0.33333333333333337)
```

$$\begin{aligned}\theta_1 &= \left(\frac{1}{2}, \frac{1}{4}, \frac{1}{4}\right) \\ \theta_2 &= \left(\frac{1}{4}, \frac{1}{2}, \frac{1}{4}\right) \\ \Pr(\mathbf{C} = \text{corpus}, \Theta = \theta_1) &= 2^{-3} \cdot 2^{-3} \cdot 2^{-1} = 2^{-7} \\ \Pr(\mathbf{C} = \text{corpus}, \Theta = \theta_2) &= 2^{-4} \cdot 2^{-3} \cdot 2^{-1} = 2^{-8} \\ \sum_{\theta \in \Theta} \Pr(\mathbf{C} = \text{corpus}, \Theta = \theta) &= 2^{-7} + 2^{-8} \\ \sum_{\theta \in \Theta} \Pr(\mathbf{C} = \text{corpus}, \Theta = \theta) &= \frac{1}{2^7} + \frac{1}{2^8} \\ \sum_{\theta \in \Theta} \Pr(\mathbf{C} = \text{corpus}, \Theta = \theta) &= \frac{2^8 + 2^7}{2^{15}} \\ \sum_{\theta \in \Theta} \Pr(\mathbf{C} = \text{corpus}, \Theta = \theta) &= \frac{2^7 \cdot (2 + 1)}{2^{15}} \\ \sum_{\theta \in \Theta} \Pr(\mathbf{C} = \text{corpus}, \Theta = \theta) &= \frac{3}{2^8} \\ \Pr(\Theta = \theta_1 | \mathbf{C} = \text{corpus}) &= \frac{2^{-7} \cdot 2^8}{3} = \frac{2^1}{3} = 0.666 \\ \Pr(\Theta = \theta_2 | \mathbf{C} = \text{corpus}) &= \frac{2^{-8} \cdot 2^8}{3} = \frac{2^0}{3} = 0.333\end{aligned}$$

The Conditional Distribution looks the way it does because of the values of θ_1 and θ_2 . Mathematically, it is proven above.

$$\begin{aligned}\Pr(\mathbf{C} = \text{corpus}, \Theta = \theta_1) &> \Pr(\mathbf{C} = \text{corpus}, \Theta = \theta_2) \\ \therefore \Pr(\Theta = \theta_1 | \mathbf{C} = \text{corpus}) &> \Pr(\Theta = \theta_2 | \mathbf{C} = \text{corpus})\end{aligned}$$

Problem 6: When you call `compute-conditional-dist`, you get back a probability distribution over values of θ (the conditional distribution over θ given an observed corpus). This is a probability distribution just like any other. In particular, it can be used as the prior distribution over values of θ in a hierarchical bag of words model. Given this new hierarchical BOW model, we can do all of the things that we normally do with such a model. In particular, we can compute the marginal likelihood of a corpus under this

model. This marginal likelihood is called a *posterior predictive distribution*.

Below we have defined the skeleton of a function `compute-posterior-predictive`. It takes three arguments: `observed-corpus`, `new-corpus`, and `theta-probs`. `observed-corpus` is a corpus which we observe, and use to compute a conditional distribution over values of θ . Given this conditional distribution over θ , we will then compute the marginal likelihood of the corpus `new-corpus`. The procedure `compute-posterior-predictive` should return the marginal likelihood of `new-corpus` given the conditional distribution on θ .

```
(defn compute-posterior-predictive [observed-corpus new-corpus theta-probs]
  (let [conditional-dist ...
        (compute-marginal ...
```

Once you have implemented `compute-posterior-predictive`, call `(compute-posterior-predictive my-corpus my-corpus theta-prior)`. What does this quantity represent? How does its value compare to the marginal likelihood that you computed in Problem 2?

Answer 6:

```
(defn compute-posterior-predictive [observed-corpus new-corpus theta-probs]
  (let [conditional-dist (exponentiate
    (compute-conditional-dist observed-corpus theta-probs))]
    (compute-marginal new-corpus conditional-dist)))
(compute-posterior-predictive my-corpus my-corpus theta-prior)
-6.2630344058337934
(Math/pow 2 (compute-posterior-predictive my-corpus my-corpus theta-prior))
0.013020833333333337
```

This is the posterior predictive of `my-corpus`. As stated in the question, the posterior predictive represents the marginal likelihood of `new-corpus` with an additional factor of the conditional distribution on θ . In this particular case, the function returns the likelihood that `new-corpus` will be the same as `observed-corpus`. The value obtained from Problem 2 is -6.415037499278844 . The value obtained from `compute-posterior-predictive` (-6.2630344058337934) is slightly greater than the value obtained in Problem 2 due to the fact stated above. Once their decimal values are calculated, it can be seen that the difference is negligible.

Problem 7: In the previous problems, we have written code that will compute marginal and conditional distributions *exactly*, by enumerating over all possible values of θ . In the next problems, we will develop an alternate approach to computing these distributions.

Instead of computing these distributions exactly, we will approximate them using random sampling.

The following functions were defined in class, and will be useful for us going forward:

```
(defn normalize [params]
  (let [sum (apply + params)]
    (map (fn [x] (/ x sum)) params)))

(defn flip [weight]
  (if (< (rand 1) weight)
      true
      false))

(defn sample-categorical [outcomes params]
  (if (flip (first params))
      (first outcomes)
      (sample-categorical (rest outcomes)
                          (normalize (rest params)))))

(defn repeat [f n]
  (if (= n 0)
      '()
      (cons (f) (repeat f (- n 1)))))

(defn sample-BOW-sentence [len probabilities]
  (if (= len 0)
      '()
      (cons (sample-categorical vocabulary probabilities)
            (sample-BOW-sentence (- len 1) probabilities))))
```

Recall that the function `sample-BOW-sentence` samples a sentence from the bag of words model of length `len`, given the parameters `theta`.

Define a function `sample-BOW-corpus`, which takes three arguments: `theta`, `sent-len`, and `corpus-len`. The argument `theta` is a value of the model parameters θ . The arguments `sent-len` and `corpus-len` are positive integers. The function should return a sample corpus from the bag of words model, given the model parameters `theta`. Each sentence should be of length `sent-len` and number of sentences in the corpus should be equal to `corpus-len`. For example, if `sent-len` equals 2 and `corpus-len` equals 2, then this function should return a list of 2 sentences, each consisting of 2 words.

Hint: Use `sample-BOW-sentence` and `repeat`.

Answer 7:

```
(defn sample-BOW-corpus [theta sent-len corpus-len]
  (repeatedly corpus-len (fn [] (sample-BOW-sentence sent-len theta))))
```

Problem 8: Below we have defined the skeleton of the function `sample-theta-corpus`. This function takes three arguments: `sent-len` `corpus-len` and `theta-probs`. It returns a list with two elements: a value of θ sampled from the distribution defined by `theta-probs`; and a corpus sampled from the bag of words model given the sampled θ . (The number of sentences in the corpus should equal `corpus-len`, and each sentence should have `sent-len` words in it.)

We will call the return value of this function a `theta-corpus` pair.

```
(defn sample-theta-corpus [sent-len corpus-len theta-probs]
  (let [theta ...
        (list theta ...
```

Answer 8:

```
(defn sample-theta-corpus [sent-len corpus-len theta-probs]
  (let [theta (sample-categorical thetas theta-probs)]
    (list theta (sample-BOW-corpus theta sent-len corpus-len))))
(sample-theta-corpus 2 2 theta1)

(defn sample-thetas-corpora [sample-size sent-len corpus-len theta-probs]
  (repeatedly sample-size
    (fn [] (sample-theta-corpus sent-len corpus-len theta-probs))))
```

Problem 9: Below we have defined some useful functions for us. The function `get-theta` takes a `theta-corpus` pair, and returns the value of `theta` in it. The function `get-corpus` takes a `theta-corpus` pair, and returns its corpus value. The function `sample-thetas-corpora` samples multiple `theta-corpus` pairs, and returns a list of them. In particular, the number of samples it returns equals `sample-size`.

```
(defn get-theta [theta-corpus]
  (first theta-corpus))

(defn get-corpus [theta-corpus]
  (first (rest theta-corpus)))
```



```
(defn sample-thetas-corpora [sample-size sent-len corpus-len theta-probs]
  (repeat (fn [] (sample-theta-corpus sent-len corpus-len theta-probs))
    sample-size))
```

We are now going to estimate the marginal likelihood of a corpus by using random sampling. Here is the general approach that we are going to use. We are going to sample some number (for example 1000) of theta-corpus pairs. These are 1000 samples from the joint distribution defined by the hierarchical bag of words model. We are then going to throw away the values of theta that we sampled; this will leave us with 1000 corpora sampled from our model.

We are going to use these 1000 sampled corpora to estimate the probability of a specific target corpus. The process here is simple. We just count the number of times that our target corpus appears in the 1000 sampled corpora. The ratio of the occurrences of the target corpus to the number of total corpora gives us an estimate of the target's probability.

More formally, let us suppose that we are given a target corpus \mathbf{t} . We will define the indicator function $\mathbb{1}_{\mathbf{t}}$ by:

$$\mathbb{1}_{\mathbf{t}}(c) = \begin{cases} 1, & \text{if } t = c \\ 0, & \text{otherwise} \end{cases} \quad (2)$$

We will sample n corpora c_1, \dots, c_n from the hierarchical bag of words model. We will estimate the marginal likelihood of the target corpus \mathbf{t} by the following formula:

$$\sum_{\theta \in \Theta} \Pr(\mathbf{C} = \mathbf{t}, \Theta = \theta) \approx \frac{1}{n} \sum_i^n \mathbb{1}_{\mathbf{t}}(c_i) \quad (3)$$

Define a procedure `estimate-corpus-marginal`, which takes five arguments: `corpus`, `sample-size`, `sent-len`, `corpus-len`, and `theta-probs`. The argument `corpus` is the target corpus whose marginal likelihood we want to estimate. `sample-size` is the number of corpora that we are going to sample from the hierarchical model (its value was 1000 in the discussion above). The arguments `corpus-len` and `sent-len` characterize the number of sentences in the corpus and the number of words in each sentence, respectively. The argument `theta-probs` is the prior probability distribution over θ for our hierarchical model.

The procedure should return an estimate of the marginal likelihood of the target corpus, using the formula defined in Equation 3.

Hint: Use `sample-thetas-corpora` to get a list of samples of theta-corpus pairs, and then use `get-corpus` to extract the corpus values from these pairs (and ignore the theta

values).

Answer 9:

```
(defn estimate-corpus-marginal [corpus sample-size sent-len
  corpus-len theta-probs]
  (/ (apply +
    (let [sample-corpora
      (sample-thetas-corpora sample-size sent-len corpus-len theta-probs)]
      (map (fn [c] (if (= c corpus) 1 0))
        (map (fn [x] (get-corpus x)) sample-corpora)))) sample-size))
```

Problem 10: Call `(estimate-corpus-marginal my-corpus 50 2 2 theta-prior)` a number of times. What do you notice? Now call `(estimate-corpus-marginal my-corpus 10000 2 2 theta-prior)` a number of times. How do these results compare to the previous ones? How do these results compare to the exact marginal likelihood that you computed in Problem 2?

Answer 10:

```
(estimate-corpus-marginal my-corpus 50 2 2 theta-prior)
1/50
(estimate-corpus-marginal my-corpus 50 2 2 theta-prior)
0
(estimate-corpus-marginal my-corpus 50 2 2 theta-prior)
3/50
(estimate-corpus-marginal my-corpus 50 2 2 theta-prior)
1/50
(estimate-corpus-marginal my-corpus 50 2 2 theta-prior)
0
(estimate-corpus-marginal my-corpus 10000 2 2 theta-prior)
109/10000
(estimate-corpus-marginal my-corpus 10000 2 2 theta-prior)
97/10000
(estimate-corpus-marginal my-corpus 10000 2 2 theta-prior)
51/5000
(estimate-corpus-marginal my-corpus 10000 2 2 theta-prior)
3/250
(estimate-corpus-marginal my-corpus 10000 2 2 theta-prior)
109/10000
```

When `(estimate-corpus-marginal my-corpus 50 2 2 theta-prior)` is called, an answer of 0 is occasionally obtained. When `sample-size` goes from 50 to 10'000, the occurrence of 0 is greatly reduced. This can be reflected with a coin-flip. $\Pr(H) = \Pr(T) = 0.500$. It is noted that the probabilities of flipping heads and flipping tails are both 0.500. If a coin was flipped 3 times, the occurrence of Heads would be 0, 0.333, 0.666 or 1.000. None of these values align exactly with 0.5. The sample size is so small that it cannot be reflective of what actually occurs.

Compared to the marginal likelihood that was obtained in Problem 2, the values obtained for `sample-size` of 10'000 are close to that of Problem 2. As stated in the previous paragraph, there was a greater error when the `sample-size` was 50.

Problem 11: These functions will be useful for us in the next problem.

```
(defn get-count [obs observation-list count]
  (if (empty? observation-list)
      count
      (if (= obs (first observation-list))
          (get-count obs (rest observation-list) (+ 1 count))
          (get-count obs (rest observation-list) count))))

(defn get-counts [outcomes observation-list]
  (let [count-obs (fn [obs] (get-count obs observation-list 0))]
    (map count-obs outcomes)))
```

In Problem 9, we introduced a way of approximating the marginal likelihood of a corpus by using random sampling. We can similarly approximate a conditional probability distribution by using random sampling.

Suppose that we have observed a corpus `c`, and we want to compute the conditional probability of a particular θ . We can approximate this conditional probability as follows. We first sample n `theta-corpus` pairs. We then remove all of the pairs in which the corpus does not match our observed corpus `c`. We finally count the number of times that θ occurs in the remaining `theta-corpus` pairs, and divide by the total number of remaining pairs. This process is called *rejection sampling*.

Define a function `rejection-sampler` which has the following form:

```
(rejection-sampler theta observed-corpus sample-size sent-len
  corpus-len theta-probs)
```

We want to get an estimate of the conditional probability of `theta`, given that we have observed the corpus `observed-corpus`. `sample-size` is a positive integer, and we will

estimate this conditional probability by taking sample-size samples from the joint distribution on theta-corpus pairs. The procedure should filter out any theta-corpus pairs in which the corpus does not equal the observed corpus. In the remaining pairs, it should then count the number of times that theta occurs, and divide by the total number of remaining pairs.

Hint: Use get-counts to count the number of occurrences of theta.

Answer 11:

```
(defn rm-2 [observed-corpus corpora-pairs]
  (filter (fn [x] (= observed-corpus (first (rest x)))) corpora-pairs))

(defn rejection-sampler [theta observed-corpus sample-size sent-len
  corpus-len theta-probs]
  (let [corpora-pairs
        (sample-thetas-corpora sample-size sent-len corpus-len theta-probs)]
    (let [matching-corpora (rm-2 observed-corpus corpora-pairs)]
      (let [sample-theta (map (fn [x] (get-theta x)) matching-corpora)]
        (if (= 0 (count matching-corpora))
            Double/NaN
            (/ (get-count theta sample-theta 0) (count matching-corpora)))))))
```

Problem 12: Call (rejection-sampler theta1 my-corpus 100 2 2 theta-prior) a number of times. What do you notice? How large does sample-size need to be until you get a stable estimate of the conditional probability of theta1? Why does it take so many samples to get a stable estimate?

Answer 12:

```
(rejection-sampler theta1 my-corpus 100 2 2 theta-prior)
1/3
(rejection-sampler theta1 my-corpus 100 2 2 theta-prior)
0
(rejection-sampler theta1 my-corpus 100 2 2 theta-prior)
NaN
(rejection-sampler theta1 my-corpus 200 2 2 theta-prior)
1/2
(rejection-sampler theta1 my-corpus 500 2 2 theta-prior)
1/3
(rejection-sampler theta1 my-corpus 1000 2 2 theta-prior)
1/5
```

```
(rejection-sampler theta1 my-corpus 5000 2 2 theta-prior)
39/53
(rejection-sampler theta1 my-corpus 10000 2 2 theta-prior)
69/122
(rejection-sampler theta1 my-corpus 15000 2 2 theta-prior)
116/171
(rejection-sampler theta1 my-corpus 15000 2 2 theta-prior)
5/8
(rejection-sampler theta1 my-corpus 15000 2 2 theta-prior)
125/174
(rejection-sampler theta1 my-corpus 17500 2 2 theta-prior)
137/201
(rejection-sampler theta1 my-corpus 17500 2 2 theta-prior)
9/14
(rejection-sampler theta1 my-corpus 17500 2 2 theta-prior)
44/63
(rejection-sampler theta1 my-corpus 20000 2 2 theta-prior)
65/98
(rejection-sampler theta1 my-corpus 20000 2 2 theta-prior)
179/276
(rejection-sampler theta1 my-corpus 20000 2 2 theta-prior)
154/239
```

When sample-size equals 100, an unreliable value is obtained. This can be explained by the phenomenon of the coin toss mentioned in Problem 10. It takes many samples in order to get a stable estimate because the more tries increase the occurrence of corpus equalling the observed corpus.

A more stable value is obtained when sample-size equals 17'500. Around this value, it starts to stabilise. When sample-size has a value of 20'000, the rejection-sampler stabilises around $\frac{2}{3}$.

In terms of accuracy, we obtain a more accurate value because we have more values to interpret and sample. When we only had 3 tries from the coin toss, we could only obtain 0, 0.333, 0.666 or 1.000. There was not a lot to sample. In terms of division and fractions with an increase in size, both our numerator and denominator increase in value to give us a more accurate value.